

Coping with NP-completeness: Approximation Algorithms

Alexander S. Kulikov

Steklov Institute of Mathematics at St. Petersburg
Russian Academy of Sciences

Advanced Algorithms and Complexity
Data Structures and Algorithms

Outline


- 1 Vertex cover
- 2 Traveling salesman
 - Metric TSP
 - Local search

Vertex cover (optimization version)


Input: A graph.

Output: A subset of vertices of minimum size that touches every edge.


Example




Example



Example




Example




ApproxVertexCover($G(V, E)$)

```
 $C \leftarrow$  empty set  
while  $E$  is not empty:  
     $\{u, v\} \leftarrow$  any edge from  $E$   
    add  $u, v$  to  $C$   
    remove from  $E$  all edges incident to  $u, v$   
return  $C$ 
```



Example




Example




Example




Example




Example




Example




Example




Example



Example



Example



Lemma

The algorithm `ApproxVertexCover` is 2-approximate: it returns a vertex cover that is at most twice as large as an optimal one and runs in polynomial time.


Proof

- The set M of all edges selected by the algorithm forms a matching
- Any vertex cover of the graph has size at least $|M|$
- The algorithm returns a vertex cover C of size $2|M|$, hence


$$|C| = 2 \cdot |M| \leq 2 \cdot \text{OPT}$$



Example



Example



Summary

- We don't know the value of OPT, but we've managed to prove that

$$|C| \leq 2 \cdot \text{OPT}$$

- This is because we know a **lower bound** on OPT: it is at least the size of any matching

$$|C| = 2 \cdot |M| \leq 2 \cdot \text{OPT}$$

Final Remarks

- The bound is tight: there are graphs for which the algorithm returns a vertex cover of size twice the minimum size.
- No 1.99-approximation algorithm is known.

Outline

- 1 Vertex cover
- 2 Traveling salesman
 - Metric TSP
 - Local search

Metric TSP (optimization version)

Input: An undirected graph $G(V, E)$ with non-negative edge weights satisfying the triangle inequality: for all $u, v, w \in V$,
$$d(u, v) + d(v, w) \geq d(u, w).$$

Output: A cycle of minimum total length visiting each vertex exactly once .

Lower Bound

- We are going to design a 2-approximation algorithm: it returns a cycle that is at most twice as long as an optimal cycle: $C \leq 2 \cdot \text{OPT}$
- Since we don't know the value of OPT , we need a good lower bound L on OPT :

$$C \leq 2 \cdot L \leq 2 \cdot \text{OPT}$$

Minimum Spanning Trees

Lemma

Let G be an undirected graph with non-negative edge weights. Then $\text{MST}(G) \leq \text{TSP}(G)$.


Proof

By removing any edge from an optimum TSP cycle one gets a spanning tree of G . \square


ApproxMetricTSP(G)

```
 $T \leftarrow$  minimum spanning tree of  $G$   
 $D \leftarrow T$  with each edge doubled  
find an Eulerian cycle  $C$  in  $D$   
return a cycle that visits vertices in  
the order of their first appearance in  $C$ 
```


Example: points on a plane




Example: points on a plane




Example: points on a plane




Example: points on a plane



Example: points on a plane



Example: points on a plane



Lemma

The algorithm `ApproxMetricTSP` is 2-approximate.

Proof

- The total length of the MST T is at most OPT .

Lemma

The algorithm `ApproxMetricTSP` is 2-approximate.

Proof

- The total length of the MST T is at most OPT .
- Bypasses can only decrease the total length.



Final Remarks

- The currently best known approximation algorithm for metric TSP is Christofides' algorithm that achieves a factor of 1.5
- If $\mathbf{P} \neq \mathbf{NP}$, then there is no α -approximation algorithm for the general version of TSP for any polynomial time computable function α

Outline

- 1 Vertex cover
- 2 Traveling salesman
 - Metric TSP
 - Local search

LocalSearch

```
s ← some initial solution
while there is a solution  $s'$  in the
neighborhood of  $s$  which is better than  $s$ :
    s ←  $s'$ 
return s
```


- Computes a local optimum instead of a global optimum
- The larger is the neighborhood, the better is the resulting solution and the higher is the running time

Local Search for TSP

- Let s and s' be two cycles visiting each vertex of the graph exactly once
- The distance between s and s' is at most d , if one can get s' by deleting d edges from s and adding other d edges
- Neighborhood $N(s, r)$ with center s and radius r : all cycles with distance at most r from s


Example

Changing two edges in a suboptimal solution:





Example

Changing two edges in a suboptimal solution:




Example

Changing two edges in a suboptimal solution:




Example

Changing two edges in a suboptimal solution:




Example

A suboptimal solution that cannot be improved by changing two edges:



Example

A suboptimal solution that cannot be improved by changing two edges:



Need to allow changing three edges to improve this solution

Performance

- Trade-off between quality and running time of a single iteration
- Still, the number of iterations may be exponential and the quality of the found cycle may be poor
- But works well in practice

Coping with NP-completeness

- special cases
- intelligent exhaustive search
- approximation algorithms